



# Austin Elastic User Group Meetup

2026 January

Srinivas Chilakapati

[elastic.co](https://elastic.co)



elastic.co | © 2026 Elasticsearch B.V. All Rights Reserved.

Follow the below steps for this tutorial.

## 1. Ingest Sample Dataset

Ingest sample dataset ([Open Food Facts data](#)) into your Elasticsearch deployment. You can use tools [this extractor](#) for data preparation.

## 2. Create a Lexical-Only Index

Create an index with standard text and keyword mappings for traditional lexical search.

```
None

{
  "openfoodinventory": {
    "aliases": {},
    "mappings": {
      "dynamic": "true",
      "properties": {
        "attr_keys": { "type": "keyword" },
        "attrs": { "type": "flattened" },
        "brand": {
          "type": "text",
          "fields": {
            "keyword": { "type": "keyword", "ignore_above": 256 }
          }
        },
        "categories": { "type": "keyword" },
        "currency": { "type": "keyword" },
        "description": { "type": "text" },
        "dietary_restrictions": { "type": "keyword" },
        "id": { "type": "keyword" },
        "image_url": { "type": "keyword", "ignore_above": 2048 },
        "nutriments": {
          "properties": {
            "energy_kcal_100g": { "type": "float" },
            "fat_g_100g": { "type": "float" },
            "fiber_g_100g": { "type": "float" },
            "protein_g_100g": { "type": "float" },
            "salt_g_100g": { "type": "float" },
            "sugars_g_100g": { "type": "float" }
          }
        }
      }
    }
  }
}
```



```

        "saturated_fat_g_100g": { "type": "float" },
        "sugars_g_100g": { "type": "float" }
    },
    "price": { "type": "scaled_float", "scaling_factor": 100
},
    "title": {
        "type": "text",
        "fields": {
            "keyword": { "type": "keyword", "ignore_above": 256 }
        }
    },
    "url": { "type": "keyword", "ignore_above": 2048 }
}
},
"settings": {
    "index": {
        "routing": {
            "allocation": {
                "include": { "_tier_preference": "data_content" }
            }
        },
        "number_of_shards": "1",
        "provided_name": "openfoodinventory",
        "default_pipeline": "openfood_nutriments_cleanup",
        "creation_date": "1768979921478",
        "number_of_replicas": "1",
        "uuid": "_ICIs_VYQJCNRrna411eIQ",
        "version": { "created": "9039003" }
    }
}
}
}

```

### Lexical Search Example:



Lexical search performs well when users know what they are searching for such as "organic gala apples" or "chocolate ice cream." or a product/manufacturer name.

```
None

# Lexical search
GET openfoodinventory/_search
{
  "_source": ["title", "description", "price"],
  "from" : 0,
  "size": "10",
  "query": {
    "bool": {
      "must": [
        {
          "multi_match": {
            "query": "give me a healthy icecream alternative",
            "fields": ["title^2", "description"],
            "slop" : "2"
          }
        }
      ]
    }
  },
  "aggs": {
    "categories": {
      "terms": {
        "field": "categories"
      }
    }
  }
}
```

However, Lexical search often fails when a user is not sure about the product they want to purchase but can describe desired product features using natural language, rather than specific product names.

- *Examples where lexical search fails:*
  - give me a healthy icecream alternative
  - healthy lunch for kids lunchbox



- find me chips that do not use seed oils

Semantic search is designed to solve this by understanding the meaning and context of the words.

### 3. Setting Up Semantic Search

Setting up semantic search is straightforward and easy with Elasticsearch:

- **3.1. Choose Your ML Model:** Select a model that generates [vector embeddings](#) to capture the semantic meaning of the text. Options include:
  - Elastic built-in models or custom models run on Elastic ML nodes.
  - [Elastic Inference Service \(EIS\)](#).
  - Third-party inference providers (OpenAI, Cohere, Huggingface, etc.).
- **3.2. Create Inference Endpoint:** Define an inference endpoint, optionally specifying a [chunking strategy](#).

None

```
PUT _inference/text_embedding/openai_embeddings
{
  "service": "openai",
  "service_settings": {
    "model_id": "text-embedding-3-small",
    "api_key": "<YOUR_OPENAI_API_KEY>",
    "similarity": "dot_product",
    "dimensions": 1536,
    "rate_limit": {
      "requests_per_minute": 3000
    },
    "chunking_settings": {
      "max_chunk_size": 300,
      "overlap": 50,
      "strategy": "word"
    }
  }
}
```



- **3.3. Identify Fields for Semantic Search:** Determine which fields (e.g., `description`) are best suited for natural language querying.
- **3.4. Create Index with `semantic_text` Fields:**  
Create a new index (`openfoodinventory_openai`) and map the chosen field (`description`) to a new `semantic_text` field (`description_embeddings`).

None

```
PUT openfoodinventory_openai
{
  "aliases": {},
  "mappings": {
    "dynamic": "true",
    "properties": {
      /* ... other field mappings remain the same ... */
      "description": {
        "type": "text",
        "copy_to": "description_embeddings"
      },
      "description_embeddings" : {
        "type": "semantic_text",
        "inference_id": "openai-text_embedding-q5xk3xshxt"
      },
      /* ... rest of the field mappings ... */
    }
  },
  "settings": {
    "index": {
      "routing": {
        "allocation": {
          "include": { "_tier_preference": "data_content" }
        }
      },
      "number_of_shards": "1",
      "number_of_replicas": "1"
    }
  }
}
```



```
}
```

The key change between lexical and semantic index is:

```
None  
"description": {  
    "type": "text",  
    "copy_to": "description_embeddings"  
},  
"description_embeddings" : {  
    "type": "semantic_text",  
    "inference_id": "openai-text_embedding-q5xk3xshxt"  
}
```

- **3.5. Ingest Data into Semantic Index:** Reindex the data. The `copy_to` and `semantic_text` mapping automatically handles the vector embedding generation.

```
None  
POST  
_reindex?wait_for_completion=false&requests_per_second=600&scroll  
=60m&slices=2  
{  
    "source": {  
        "index": "openfoodinventory",  
        "size": 500  
    },  
    "dest": {  
        "index": "openfoodinventory_openai"  
    }  
}
```

### Semantic/Vector Search Example:

Query the `description_embeddings` field using the user's natural language query.



```
None

GET openfoodinventory_openai/_search
{
  "_source": ["title", "description"],
  "query": {
    "match": {
      "description_embeddings": {
        "query": "give me a healthy icecream alternative"
      }
    }
  }
}
```

### Hybrid Search Example (Best of Both Worlds):

Combine the precision of lexical search with the relevance of semantic search using the [retriever API](#).

```
None

GET openfoodinventory_openai/_search
{
  "from": 0,
  "size": 10,
  "_source": ["title", "description", "price"],
  "retriever": {
    "linear": {
      "retrievers": [
        {
          "retriever": {
            "standard": {
              "query": {
                "knn": {
                  "field": "description_embeddings",
                  "query_vector_builder": {
                    "text_embedding": {
                      "model_id": "openai-text_embedding-q5xk3xshxt",

```



```
        "model_text": "best olive oil, under $10"
    }
},
"k": 10,
"num_candidates": 100
}
}
},
"weight": "1"
},
{
"retriever": {
"standard": {
"query": {
"multi_match": {
"query": "best value olive oil, under $10",
"fields": [
"title^2",
"description"
],
"slop": "2"
}
}
}
},
"weight": "1"
}
],
"normalizer": "minmax"
}
}
}
```

This hybrid query successfully returns more relevant results than a lexical-only search.



## Limitations of Vector Search:

Semantic search fails in below examples because the model focuses its attention on the context and meaning of the text rather than the specifics/attributes of the request like 5g sugar or price < 10\$.

- *Examples where vector search fails:*

- breakfast cereal with less than 5g sugar per serving
- snack bar with no more than 6 ingredients
- healthy ice cream under \$10

These attributes must be extracted from the user's query and applied as filters to the search. This is where **Elastic Agent Builder** can be leveraged to create an agent that extracts these attributes and generates the final, filtered query.

The screenshot shows the Elastic Agent Builder interface with the following details:

- Tool ID:** openfoodinventory\_hybrid\_search
- Description:** Always query openfoodinventory\_openai index  
Add filters as needed when the user specifies constraints, e.g.:  
Price limit (e.g., "under \$10")  
nutrients (e.g., "5g sugar"), etc.  
## Attribute extraction rules (map user language → fields)  
### Price → 'price'



The screenshot shows the Elastic AI Assistant interface. At the top, there are navigation icons for Home, Deployments, Agents, and a search bar. The main title is "Agents" and the sub-page title is "ecommerce\_search\_assistant". On the left, there's a sidebar with "Settings" and "Tools" tabs, where "Tools" has a count of 4. Below this is a section titled "System references" with a brief description: "Used behind the scenes to identify and guide the agent's behavior. Not shown to end users." To the right, under "Agent ID", the value "ecommerce\_search\_assistant" is listed. In the "Custom Instructions" section, there is a WYSIWYG editor toolbar and a "Preview" button. The preview content is as follows:

```
## Search Interpretation & Execution Agent (Elastic AI Agent)

You are a **search interpretation and execution agent** for an e-commerce catalog. Your purpose is to transform a user's free-text shopping query into the best product search results by combining **intent parsing**, **personalization from order history**, and **hybrid product retrieval**.

### Core duties (always follow in order)
1. **Interpret the user's query**
   - Identify the primary product intent (e.g., "olive oil", "protein bar", "low sugar cereal").
```

