Connect AI to the web with tavily

Nebius to acquire Tavily

NEBIUS | tavily

Building the unified agent stack

How does Tavily fit into my agent stack?

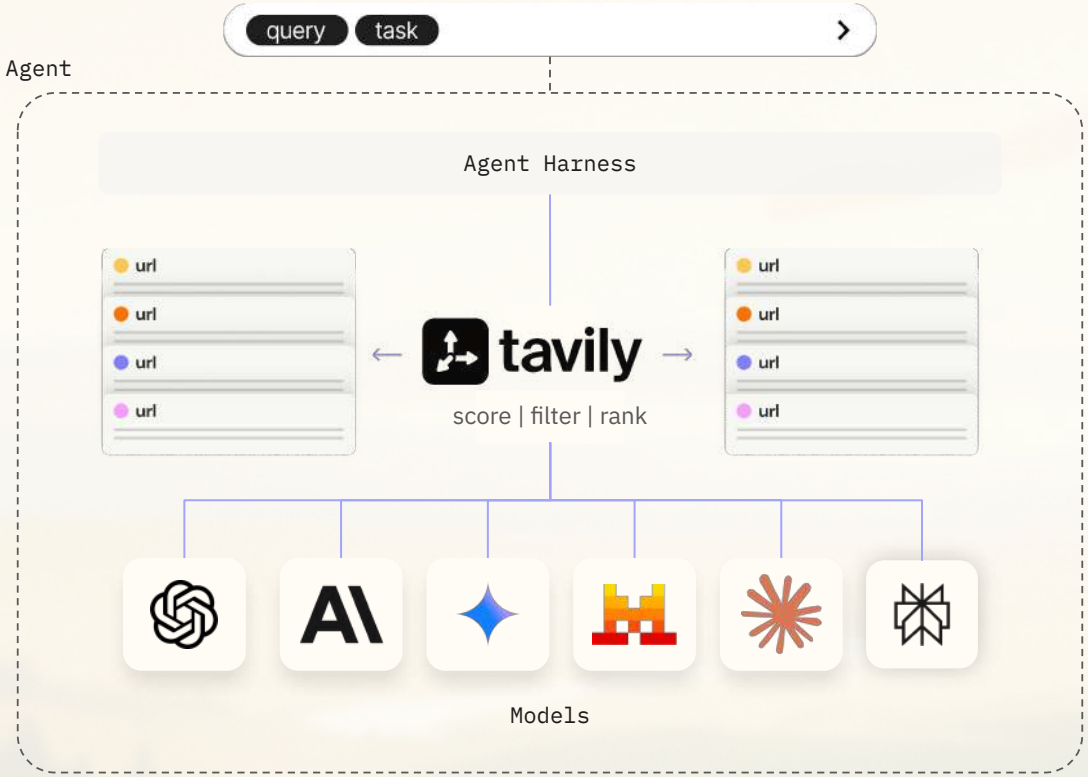

tavily

So how do we interact with the web?

Everything your agent needs to read and reason over the web

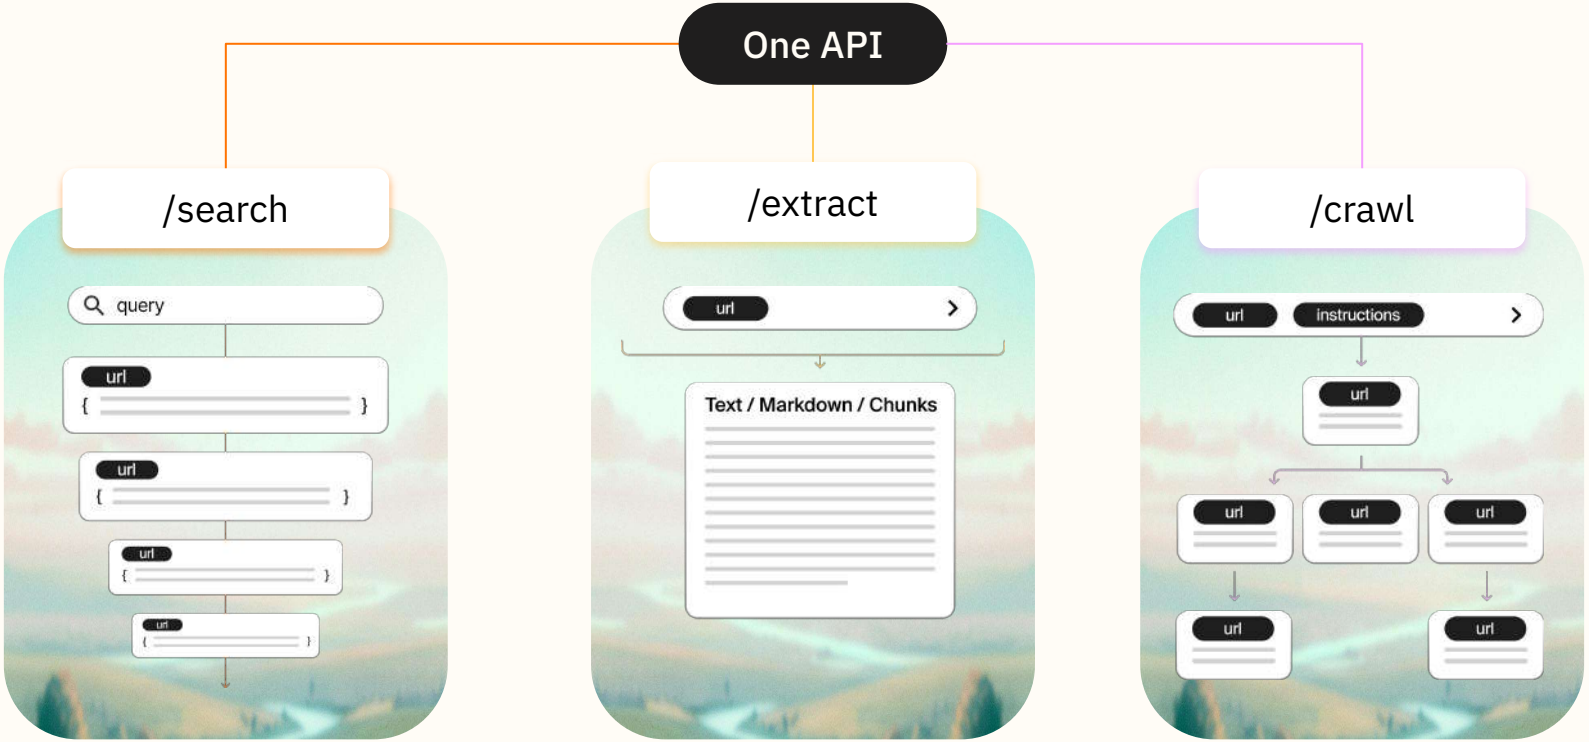

/research

Production-grade deep research agent. Delivered in one API call.

State of the art on DeepResearch Bench

You provide:

- Research query
- Operational configuration: model, schema, citation format, response type

/research agent:

- Interprets research query
- Expands task into subtasks
- Executes concurrent searches & retrieval
- Synthesizes findings across sources

You receive:

- Structured, citation-backed report

Customer Success Agent

Web Search + RAG

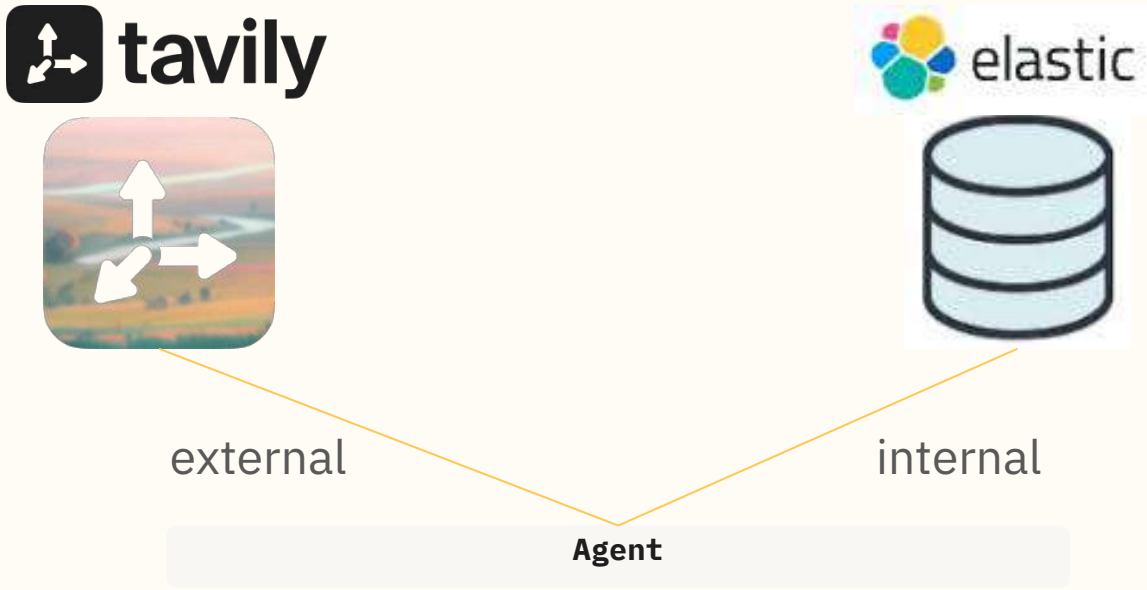

Web Search + RAG

- Recent news
- Customers shifts and colab
- New products etc

- Communication
- Meetings recaps
- Usage

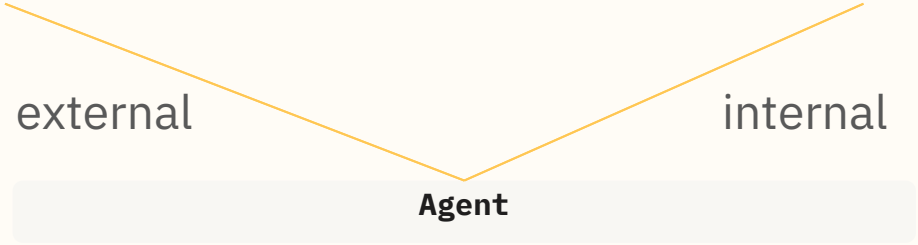

Agent Architecture

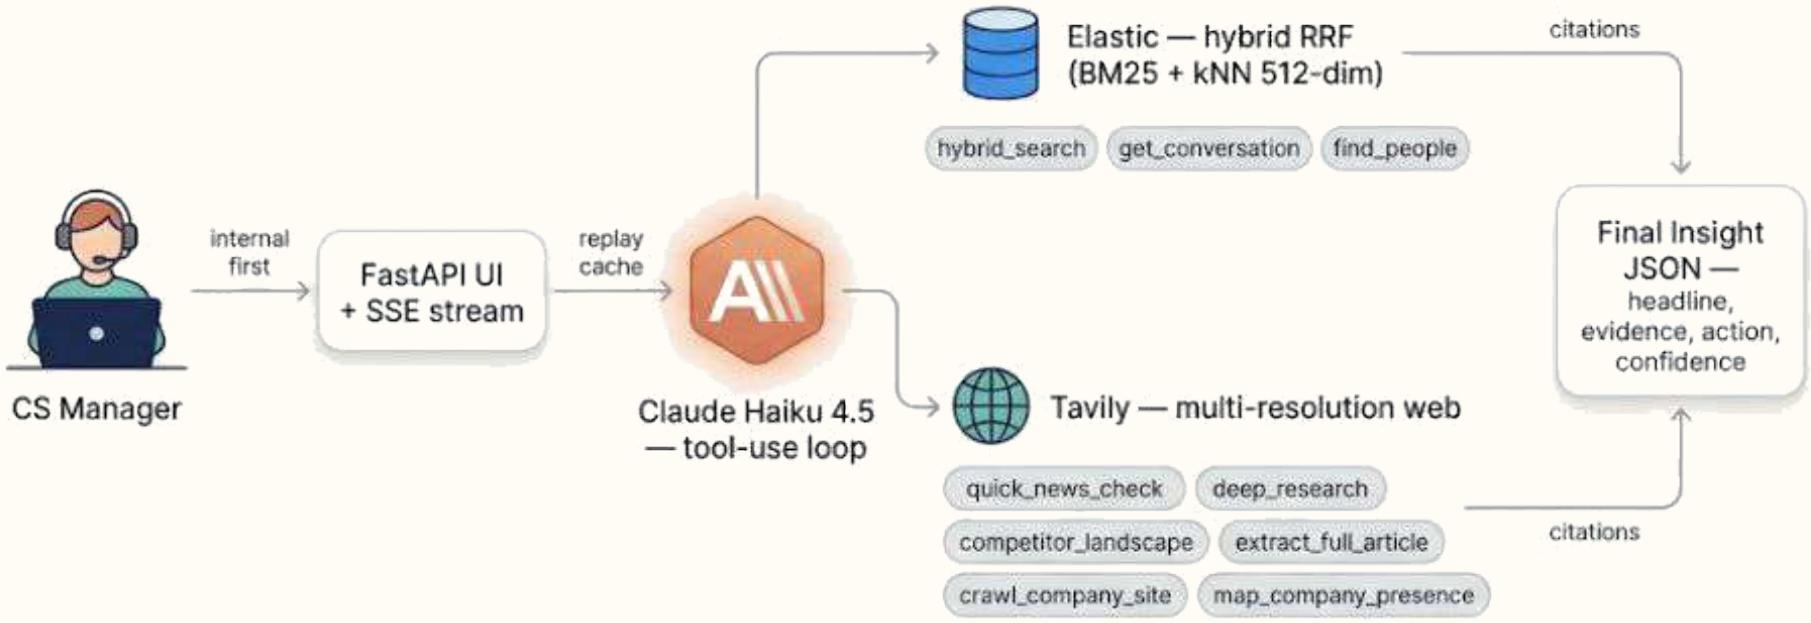

RAG - Architecture

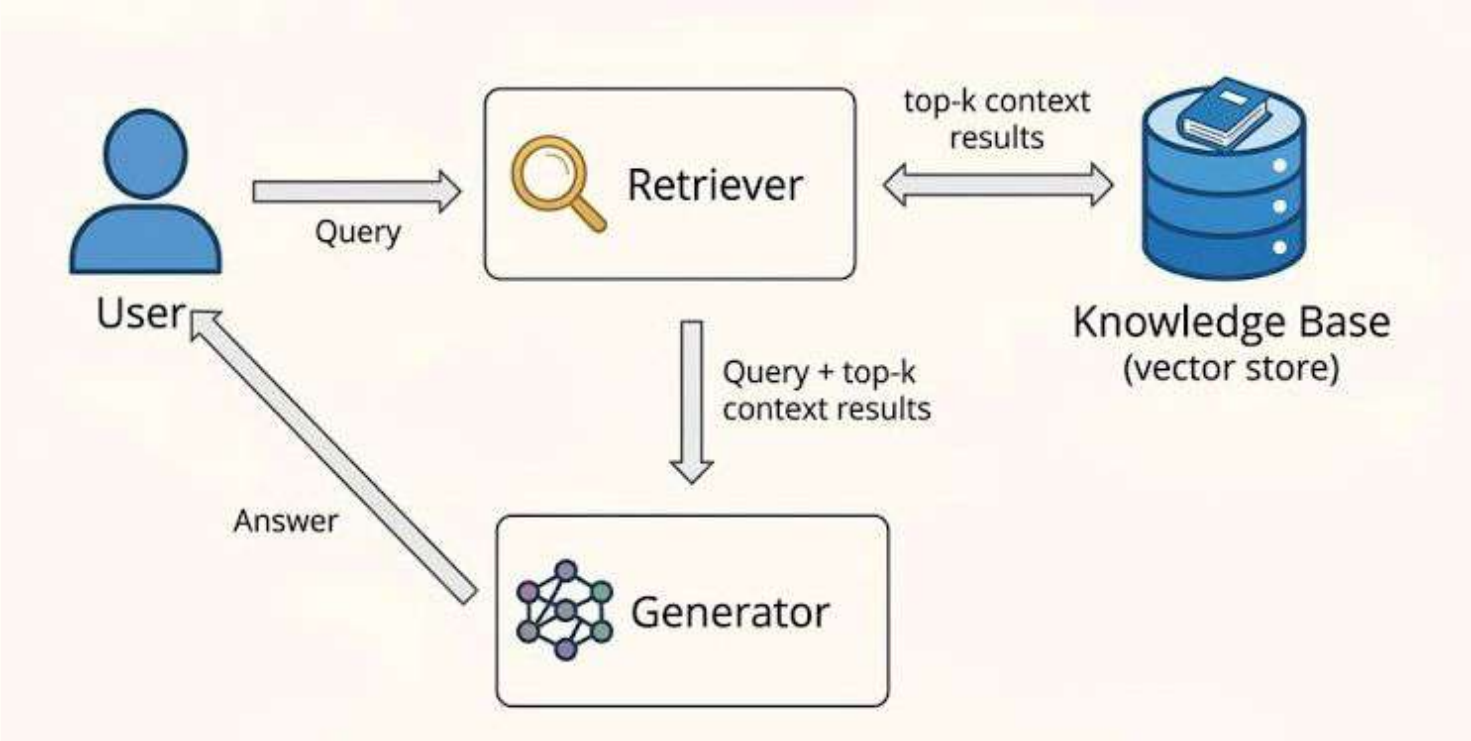

Elastic Usage

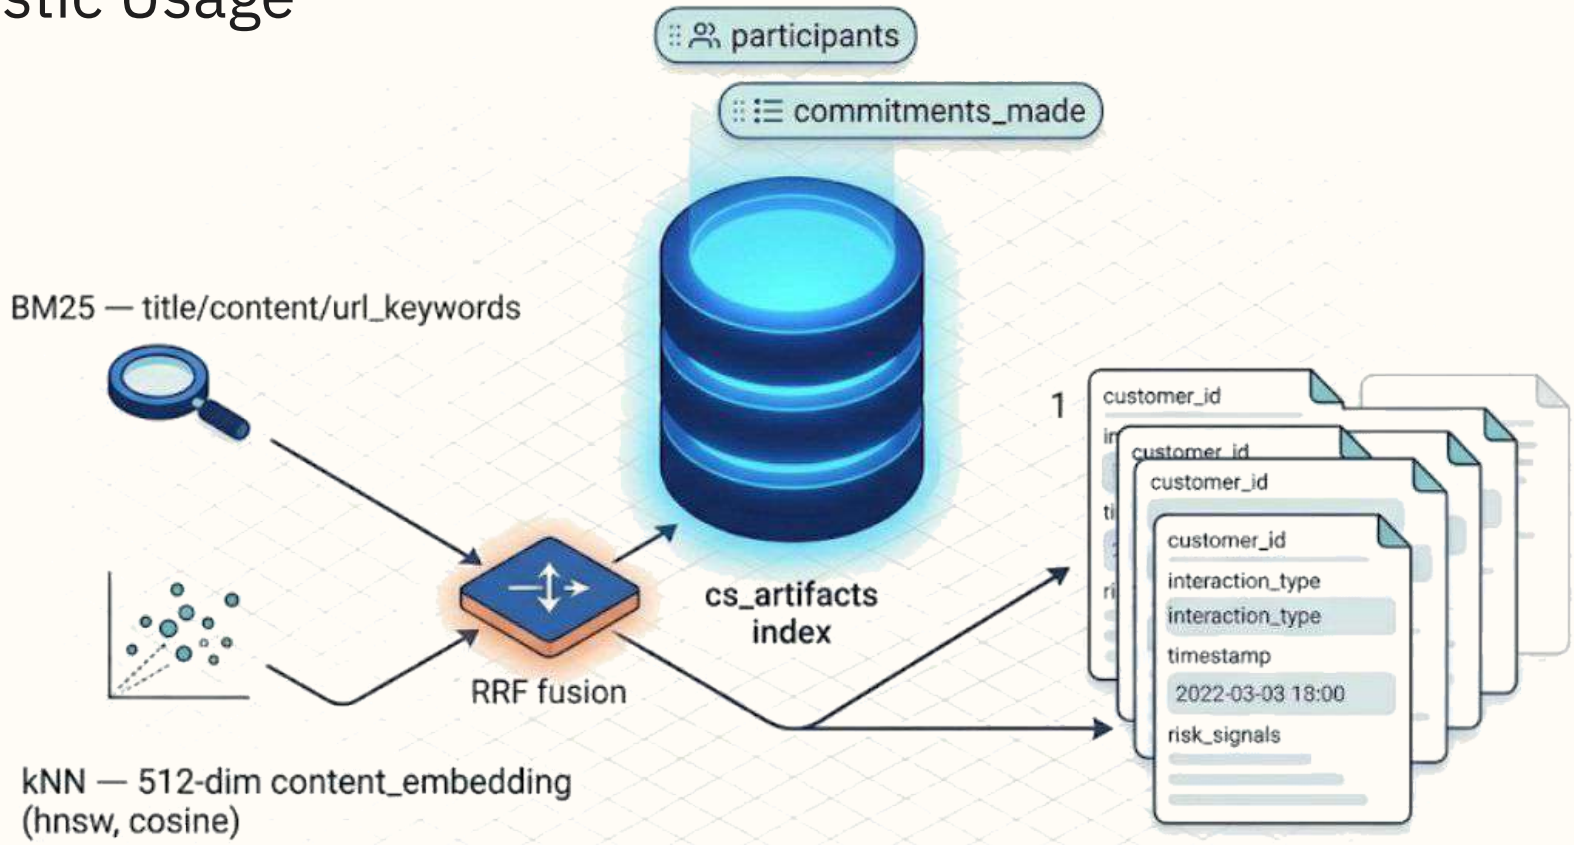

Tavily Usage

cheap & shallow — multi-resolution — slow & deep

1)

1) quick_news_check

Rapidly check for latest news on a topic or entity.

~2s

2)

2) deep_research

Comprehensive, in-depth **research** with synthesis.

30–90s

3)

3) competitor_landscape

Analyze and map competitor ecosystem.

30–60s

4)

4) extract_full_article

Extract the complete text from a full article.

~3s

5)

5) crawl_company_site

Crawl and analyze a specified company website.

10–30s

6)

6) map_company_presence

Map a company's global physical and digital presence.

5–15s

replay cache — demo/tavily_cache/

Start Building For Free

tavily