


Is Your AI Lying to you? with Live Demo

Prepared by: Michael Calizo
Principal Customer Architect - APJ, Elastic
Prepared for: Manila Elastic Meetup



elastic

| The Search
AI Company



Is Your AI Lying to You?

It's not lying on purpose, it's predicting, not fact-checking.

Here's what grounding actually fixes, proven on live **ATO** and **AU** super rules.

Elastic Agent Builder

Jina Reader

Claude

Agent Evals



Grounded, not guessed

Sourced, checkable, and fixable answers.

- ✓ Explicitly linked to verifiable source data
- ✓ Eliminates confident hallucinations
- ✓ Allows immediate peer-review correction



It's Not Malice. It's Missing Grounding.

An LLM has no built-in fact-checker. It can't tell you "I'm not sure this is still current" unless something forces it to check.



Predicts patterns, not truth

It generates the most statistically likely next words – never checks them against reality.



No sense of right or wrong

There's no internal signal that flags "I'm guessing" versus "I know this for certain."



Frozen the day training ends

Real-world rules, like super caps, keep changing long after that freeze date.

The result: fluent, confident answers that are sometimes just wrong – *functionally indistinguishable from a lie, even without intent.*



The Fix: Ground It in Something Real

Five steps turn a live webpage into a citable, tool-checked answer giving the model something to check itself against.




Same pipeline scans, indexes, and grounds any topic: **add a URL to SOURCES and re-run.**



Same Question, Different Pathways

Comparing memory-based LLM generation against a grounded agentic workflow.



Live Demo Part 2

01

ATO Tax
Question

02

End to End
Flow

03

Datasets
Review

04

Evaluation
(AI as
Judge)

05

LLM Cost
(Anthropic
Metrics)



Four Ways, Side by Side



Ungrounded LLM

- No tools or retrieval
- Answers from memory
- No source, no year



Grounded Agent

- Tools + grounding skill
- Cites source + scan date
- Admits topic gaps



Stale Agent

- Same tools + skill
- Grounded on FY21-22 data
- Cites real, outdated source



AI as Judge

- Scores tool use
- Scores answer accuracy
- Pass/fail, not vibes

Grounding doesn't stop AI from being wrong. It stops “wrong” from being silent, turning a confident guess into a sourced, checkable, fixable answer.

