



Observability and SLOs with Elastic

Date: January 29, 2026

Observability and SLOs with the Elastic Stack

**Building Reliable Systems with Service Level
Objectives**

Rajesh Kesavalalji | 01/29/2026 | Elastic Seattle
User Group

AGENDA

1. Introduction to Observability and SLOs
2. Understanding SLIs, SLOs, and Error Budgets
3. Elastic Stack for Observability
4. Demo: Checkout API with SLO Tracking
5. Key Takeaways

WHAT IS OBSERVABILITY?

The Three Pillars



Metrics: Quantitative measurements over time

- Request rate, latency, error rate, resource utilization



Logs: Discrete events with timestamps

- Application logs, access logs, error logs




Traces: Request flow through distributed systems

- Distributed tracing, span context, service dependencies

Observability = Understanding system behavior from the outside

WHY OBSERVABILITY MATTERS



From Reactive to Proactive

✗ **Traditional Monitoring:** "Something is broken, let's investigate"

✓ **Observability:** "Why is the system behaving this way?"



Benefits with Observability

- ⚡ Faster incident resolution
- 🔮 Proactive problem detection
- 😊 Better user experience
- 📈 Data-driven decisions

SERVICE LEVEL OBJECTIVES (SLOs)

What are SLOs?

SLO: A target level of reliability for a service

Example: "99.9% of requests should succeed"

Purpose:

- Set clear reliability expectations
- Enable data-driven decisions
- Balance reliability vs. feature velocity

SLO ≠ SLA

- SLO: Internal target
- SLA: External commitment (with consequences)

SERVICE LEVEL INDICATORS (SLIs)

What We Measure

SLI: A quantitative measure of service reliability

Common SLIs:

-  **Availability:** Percentage of successful requests
-  **Latency:** Response time (p50, p95, p99)
-  **Throughput:** Requests per second
-  **Error Rate:** Percentage of failed requests

Example: "99.9% of checkout requests succeed within 200ms"

ERROR BUDGETS

The Safety Margin

Error Budget: The acceptable amount of unreliability

Calculation: $100\% - \text{SLO Target}$

Example:

- SLO: 99.9% availability
- Error Budget: 0.1% (43.2 minutes/month)

Purpose:

- Enable risk-taking
- Guide deployment decisions
- Balance reliability vs. innovation

BURN-RATE ALERTS

Early Warning System

Burn Rate: How fast error budget is consumed



Fast-Burn Alert: Error budget consumed 2x faster than sustainable

- Short window (6 hours)
- Immediate action required



Slow-Burn Alert: Error budget consumed 1.5x faster than sustainable

- Long window (30 days)
- Proactive intervention

Goal: Catch issues before SLO violation



ELASTIC STACK OVERVIEW

The Observability Platform



Elasticsearch: Search and analytics engine

- Stores metrics, logs, traces
- Real-time indexing and querying



Kibana: Visualization and dashboards

- APM UI, SLO tracking, custom dashboards



APM Server: APM data collection


- Receives data from APM agents
- Processes and indexes to Elasticsearch



APM Agents: Application instrumentation

- Automatic transaction and span tracking

DEMO ARCHITECTURE



Checkout API with Elastic Stack

DEMO: CHECKOUT API

What We're Building

Service: E-commerce checkout API

SLIs:

✓ Availability: 99.5% success rate

⌚ Latency: p95 < 200ms

Instrumentation: Elastic APM Node.js agent

Scenarios:

- Normal operation
- High latency (database issues)
- High error rate (payment failures)

Github:

<https://github.com/rajesharma470/elastic-observability-demo>



DEMO FLOW - PART 1

Setting Up SLOs


1. Create Availability SLO:

- Service: `checkout-api`
- Transaction: `POST /checkout`
- Target: 99.50%
- Window: 30 days

2. Create Latency SLO:

- Service: `checkout-api`
 - Transaction: `POST /checkout`
 - Target: p95 < 200ms (for demo 99.50%)
 - Window: 30 days
3. Baseline: Normal operation with 0% errors

SLOs



DEMO FLOW - PART 2

Simulating Issues


1. Latency Scenario:

- Simulate slow database (5x latency multiplier)
- Watch p95 latency exceed 200ms
- Observe latency SLO violation
- Error budget consumption

2. Error Scenario:

- Simulate payment failures (10% error rate)
- Watch availability drop below 99.50%
- Observe availability SLO violation
- Error budget burn

SLOs



KEY FEATURES DEMONSTRATED

What We Saw

- ✓ Automatic Instrumentation: No code changes needed for basic APM
- ✓ Real-time Monitoring: Live metrics and traces
- ✓ SLO Tracking: Built-in SLO management in Kibana
- ✓ Error Budget Tracking: Visual error budget consumption
- ✓ Burn-Rate Alerts: Proactive alerting

BEST PRACTICES

SLO Implementation Tips

1. **Start Simple:** Begin with 1-2 critical SLIs
2. **User-Focused:** Measure what users experience
3. **Realistic Targets:** Set achievable SLOs
4. **Review Regularly:** Adjust based on data
5. **Document Decisions:** Why this SLO? Why this target?
6. **Error Budgets:** Use them to guide deployments
7. **Automation:** Automate SLO tracking and alerts

Thank you

LinkedIn: <https://www.linkedin.com/in/rajesh-sharma-sfbay/>

Github Repo: <https://github.com/rajesharma470/elastic-observability-demo>