

Elastic multimodal search

Text, image, and video — one model, one index

September 2026

Two demos

01

Multimodal Artwork Search

65k artworks · 68k images · hybrid search · agentic chat

02

Video Moment Search

8-second chunks · 4-second overlap · timestamp-level retrieval

Search anything you can see or describe

National Gallery of Art — 65k artworks, 68k images — with rich provenance, lineage, and art-historical metadata

Hybrid Search

BM25 + vector in one query. Keyword precision for artist names and accession IDs alongside semantic similarity for "impressionist landscapes at dusk." No separate pipelines.

Filtering & Faceting

Narrow across medium, period, classification, and artist without re-embedding. Sub-second across 133k documents — exact, fast, no index rebuild.

Provenance & Ontology

Each artwork carries structured metadata: acquisition history, prior ownership, artist biography, creation context, art movement, and subject classification. Not just image vectors — a real knowledge graph, indexed and queryable.

Agentic Chat

Agent Builder wires an LLM to the index — and because the metadata is rich, the agent's answers are too. "Why are these two works related?" gets a grounded, citable answer: shared provenance, the same collector, the same period of an artist's career — not a vague similarity score.

Rich metadata powers agentic discovery



NGA collection in vector space

COLLECTION

65,000 artworks · 68,000 images



Jina v5 Omni · 1,024 dims → 3D UMAP

ART MOVEMENTS

- Dutch Golden Age
- Impressionism
- Baroque
- Renaissance
- Modern
- Photography

INTERACT

- drag rotate
- right-drag pan
- scroll zoom
- hover inspect artwork




From artwork to vector embedding

VECTORIZATION PIPELINE

1,024-DIM SPACE — 3D UMAP PROJECTION

36 NGA ARTWORKS — 6 ART MOVEMENTS



⚡ Vectorize with Jina v5 Omni



DRAG TO ROTATE • SCROLL TO ZOOM • HOVER TO INSPECT

DEMO 01

Multimodal Artwork Search

Search by text, upload an image, or ask the agent — all against the same
133k-document index.



Find the **exact moment** that matters

Natural language → timestamp — no scrubbing, no manual tagging

8-second chunks · 4-second overlap

Sub-scene granularity. The 4-second overlap ensures no relevant moment falls through a hard cut — every transition is covered by at least two chunks.

Same model, same index

Jina v5 Omni embeds video frames and text queries into the same vector space. One query type, one index — no separate video pipeline, no modality mismatch.

Timestamp-level results

Each chunk carries its start and end timecode. Search results link directly to the exact playback position — find the moment, not just the file.

Zero manual annotation


No human tagging, no caption pipeline. The embedding does the semantic lifting across the full video corpus — ingest and index, then search.

How **video moment search** is built



Breaking video into searchable moments

VIDEO TIMELINE




Clip window **8 seconds**

Step / overlap **4s · 50% shared**

CLIP 4 / 80 12s - 20s

VISUAL · CLIP FRAMES



scene: Pippin flees fell beast attack — Minas Tirith courtyard

Jina v5 Omni → 1,024-dim visual vector

AUDIO · TRANSCRIPT



"The City is under attack! Sound the alarm!" — Pippin, stone courtyard, orchestral swell

ELASTICSEARCH DOCUMENT

```
{
  "clip_id": "moment_0003",
  "video_id": "56dcc71a-2b6",
  "start_time": 12,
  "end_time": 20,
  "scene": "Pippin flees fell beast attack — Minas T...",
  "visual_vec": [-0.456, 0.428, -0.005, -0.779, ... +1020, ... +1020],
  "audio_vec": [-0.412, 0.768, 0.052, -0.559, ... +1020, ... +1020],
  "caption": "Pippin flees fell beast attack"
}
```

DEMO 02

Video Moment Search

Type a natural language query — get the exact timestamp. No scrubbing,
no chapters, no manual tagging.



● ANNOUNCING TOMORROW

jina-ocr-v1



Document images in, clean Markdown out. PDFs, slides, forms, tables — whatever your pipeline throws at it.



Built for indexing. Solves the hard problem of getting rich visual documents into Elasticsearch in the first place.




Weights are ours. Weights are yours. Trained by Elastic — fully downloadable, no cloud lock-in.

FULL BLOG · BENCHMARKS · EXAMPLES

Follow Elastic Search Labs — dropping tomorrow

JINA-OCR-V1 · DOCUMENT UNDERSTANDING

INPUT — DOCUMENT IMAGE



jina-ocr-v1 >

OUTPUT — CLEAN MARKDOWN

Quarterly Results

Overview

Revenue grew 24% YoY driven by...

 Regional breakdown

Region	Q3	Q4
--------	----	----

-----	----	----
-------	------	------

EMEA	41%	44%
------	-----	-----

AMER	38%	39%
------	-----	-----